

aml xgboost

December 8, 2023

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.metrics import classification_report, confusion_matrix
from xgboost import XGBClassifier
from scipy.stats import uniform, randint
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from joblib import dump

# Define the data types for each column (replace this with your actual dtypes_
↳dictionary)
dtypes = {
    'MachineIdentifier': 'category',
    'ProductName': 'category',
    'EngineVersion': 'category',
    'AppVersion': 'category',
    'AvSigVersion': 'category',
    'IsBeta': 'int8',
    'RtpStateBitfield': 'float16',
    'IsSxsPassiveMode': 'int8',
    'DefaultBrowsersIdentifier': 'float16',
    'AVProductStatesIdentifier': 'float32',
    'AVProductsInstalled': 'float16',
    'AVProductsEnabled': 'float16',
    'HasTpm': 'int8',
    'CountryIdentifier': 'int16',
    'CityIdentifier': 'float32',
    'OrganizationIdentifier': 'float16',
    'GeoNameIdentifier': 'float16',
    'LocaleEnglishNameIdentifier': 'int8',
    'Platform': 'category',
    'Processor': 'category',
    'OsVer': 'category',
    'OsBuild': 'int16',
    'OsSuite': 'int16',
```

```

'OsPlatformSubRelease': 'category',
'OsBuildLab': 'category',
'SkuEdition': 'category',
'IsProtected': 'float16',
'AutoSampleOptIn': 'int8',
'PuaMode': 'category',
'SMode': 'float16',
'IeVerIdentifier': 'float16',
'SmartScreen': 'category',
'Firewall': 'float16',
'UacLuaenable': 'float32',
'Census_MDC2FormFactor': 'category',
'Census_DeviceFamily': 'category',
'Census_OEMNameIdentifier': 'float16',
'Census_OEMModelIdentifier': 'float32',
'Census_ProcessorCoreCount': 'float16',
'Census_ProcessorManufacturerIdentifier': 'float16',
'Census_ProcessorModelIdentifier': 'float16',
'Census_ProcessorClass': 'category',
'Census_PrimaryDiskTotalCapacity': 'float32',
'Census_PrimaryDiskTypeName': 'category',
'Census_SystemVolumeTotalCapacity': 'float32',
'Census_HasOpticalDiskDrive': 'int8',
'Census_TotalPhysicalRAM': 'float32',
'Census_ChassisTypeName': 'category',
'Census_InternalPrimaryDiagonalDisplaySizeInInches': 'float16',
'Census_InternalPrimaryDisplayResolutionHorizontal': 'float16',
'Census_InternalPrimaryDisplayResolutionVertical': 'float16',
'Census_PowerPlatformRoleName': 'category',
'Census_InternalBatteryType': 'category',
'Census_InternalBatteryNumberOfCharges': 'float32',
'Census_OSVersion': 'category',
'Census_OSArchitecture': 'category',
'Census_OSBranch': 'category',
'Census_OSBuildNumber': 'int16',
'Census_OSBuildRevision': 'int32',
'Census_OSEdition': 'category',
'Census_OSSkuName': 'category',
'Census_OSInstallTypeName': 'category',
'Census_OSInstallLanguageIdentifier': 'float16',
'Census_OSUILocaleIdentifier': 'int16',
'Census_OSWUAutoUpdateOptionsName': 'category',
'Census_IsPortableOperatingSystem': 'int8',
'Census_GenuineStateName': 'category',
'Census_ActivationChannel': 'category',
'Census_IsFlightingInternal': 'float16',
'Census_IsFlightsDisabled': 'float16',

```

```

        'Census_FlightRing': 'category',
        'Census_ThresholdOptIn': 'float16',
        'Census_FirmwareManufacturerIdentifier': 'float16',
        'Census_FirmwareVersionIdentifier': 'float32',
        'Census_IsSecureBootEnabled': 'int8',
        'Census_IsWIMBootEnabled': 'float16',
        'Census_IsVirtualDevice': 'float16',
        'Census_IsTouchEnabled': 'int8',
        'Census_IsPenCapable': 'int8',
        'Census_IsAlwaysOnAlwaysConnectedCapable': 'float16',
        'Wdft_IsGamer': 'float16',
        'Wdft_RegionIdentifier': 'float16',
        'HasDetections': 'int8'
    }

# Define the true numerical columns
true_numerical_columns = [
    'Census_ProcessorCoreCount',
    'Census_PrimaryDiskTotalCapacity',
    'Census_SystemVolumeTotalCapacity',
    'Census_TotalPhysicalRAM',
    'Census_InternalPrimaryDiagonalDisplaySizeInInches',
    'Census_InternalPrimaryDisplayResolutionHorizontal',
    'Census_InternalPrimaryDisplayResolutionVertical',
    'Census_InternalBatteryNumberOfCharges'
]

# Load the data with specific data types
train = pd.read_csv('train.csv', dtype=dtypes)
train = train.sample(frac=0.20, random_state=42)

# Drop the columns as specified earlier
cols_to_drop = [
    'MachineIdentifier', 'AvSigVersion', 'IsBeta', 'RtpStateBitfield',
    'DefaultBrowsersIdentifier', 'AVProductsInstalled', 'HasTpm',
    'LocaleEnglishNameIdentifier', 'PuaMode', 'SMode', 'SmartScreen',
    'UacLuaenable', 'Census_DeviceFamily', 'Census_ProcessorClass',
    'Census_InternalBatteryType', 'Census_IsPortableOperatingSystem',
    'Census_IsFlightingInternal', 'Census_ThresholdOptIn',
    'OrganizationIdentifier', 'Census_IsWIMBootEnabled'
]
train.drop(columns=cols_to_drop, inplace=True)

# Separate the features and target variable
X = train.drop('HasDetections', axis=1)
y = train['HasDetections']

# Define numerical and categorical columns after dropping the unwanted columns

```

```

numerical_columns = [col for col in true_numerical_columns if col in X.columns]
categorical_columns = [col for col in X.columns if col not in numerical_columns]

# Preprocessing for numerical data: imputing missing values with median
numerical_transformer = SimpleImputer(strategy='median')

# Preprocessing for categorical data: imputing missing values with the most
# frequent value then applying OneHotEncoder
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

# Bundle preprocessing for numerical and categorical data
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numerical_transformer, numerical_columns),
        ('cat', categorical_transformer, categorical_columns)
    ])

```

```

[2]: # Define the model with added parameters to suppress warnings
model = XGBClassifier(random_state=42, use_label_encoder=False,
    eval_metric='logloss')

# Bundle preprocessing and modeling code in a pipeline
clf = Pipeline(steps=[('preprocessor', preprocessor),
    ('model', model)
])

# Split the data
X_dev, X_test, y_dev, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_dev, y_dev, test_size=0.25,
    random_state=42)

# Train the model
clf.fit(X_train, y_train)

# Predict and evaluate the model
y_train_pred = clf.predict(X_train)
y_val_pred = clf.predict(X_val)

print("Classification report for training data:")
print(classification_report(y_train, y_train_pred))

print("Classification report for validation data:")
print(classification_report(y_val, y_val_pred))

```

```

print("Confusion Matrix for training data:")
print(confusion_matrix(y_train, y_train_pred))

print("Confusion Matrix for validation data:")
print(confusion_matrix(y_val, y_val_pred))

```

Classification report for training data:

	precision	recall	f1-score	support
0	0.65	0.60	0.63	535420
1	0.63	0.67	0.65	535157
accuracy			0.64	1070577
macro avg	0.64	0.64	0.64	1070577
weighted avg	0.64	0.64	0.64	1070577

Classification report for validation data:

	precision	recall	f1-score	support
0	0.64	0.60	0.62	178629
1	0.62	0.66	0.64	178231
accuracy			0.63	356860
macro avg	0.63	0.63	0.63	356860
weighted avg	0.63	0.63	0.63	356860

Confusion Matrix for training data:

```

[[323891 211529]
 [175852 359305]]

```

Confusion Matrix for validation data:

```

[[106426  72203]
 [ 60129 118102]]

```

```

[3]: # Hyperparameter tuning with RandomizedSearchCV
parameters = {
    'model__max_depth': randint(3, 10),
    'model__learning_rate': uniform(0.01, 0.6),
    'model__n_estimators': randint(100, 1000),
}

random_search = RandomizedSearchCV(clf, param_distributions=parameters,
    n_iter=10, scoring='roc_auc', cv=5,
    n_jobs=-1, random_state=42)
random_search.fit(X_train, y_train)

# Save the best model

```

```

best_model = random_search.best_estimator_
dump(best_model, 'best_model.joblib')

# Evaluating the best model
y_train_best_pred = best_model.predict(X_train)
y_val_best_pred = best_model.predict(X_val)

print("Classification report for training data with the best model:")
print(classification_report(y_train, y_train_best_pred))

print("Classification report for validation data with the best model:")
print(classification_report(y_val, y_val_best_pred))

print("Confusion Matrix for training data with the best model:")
print(confusion_matrix(y_train, y_train_best_pred))

print("Confusion Matrix for validation data with the best model:")
print(confusion_matrix(y_val, y_val_best_pred))

```

Classification report for training data with the best model:

	precision	recall	f1-score	support
0	0.67	0.63	0.65	535420
1	0.65	0.69	0.67	535157
accuracy			0.66	1070577
macro avg	0.66	0.66	0.66	1070577
weighted avg	0.66	0.66	0.66	1070577

Classification report for validation data with the best model:

	precision	recall	f1-score	support
0	0.65	0.60	0.62	178629
1	0.63	0.67	0.65	178231
accuracy			0.64	356860
macro avg	0.64	0.64	0.64	356860
weighted avg	0.64	0.64	0.64	356860

Confusion Matrix for training data with the best model:

```

[[335506 199914]
 [164995 370162]]

```

Confusion Matrix for validation data with the best model:

```

[[107940 70689]
 [ 59242 118989]]

```

```
[8]: print(random_search.best_params_)
```

```
{'model__learning_rate': 0.2347240713084175, 'model__max_depth': 7,
'model__n_estimators': 370}
```

```
[4]: import numpy as np

# Predict on the test set using the best model
y_test_pred = best_model.predict(X_test)

# Evaluate the model's performance on the test set
test_accuracy = np.mean(y_test_pred == y_test)
print("Test Accuracy:", test_accuracy)

# Optional: Detailed classification report and confusion matrix for test data
print("Classification report for test data:")
print(classification_report(y_test, y_test_pred))

print("Confusion Matrix for test data:")
print(confusion_matrix(y_test, y_test_pred))
```

Test Accuracy: 0.6368015468250855

Classification report for test data:

	precision	recall	f1-score	support
0	0.65	0.61	0.63	179012
1	0.63	0.67	0.65	177848
accuracy			0.64	356860
macro avg	0.64	0.64	0.64	356860
weighted avg	0.64	0.64	0.64	356860

Confusion Matrix for test data:

```
[[108677  70335]
 [ 59276 118572]]
```

```
[14]: # Extracting feature names after OneHotEncoder
# Extracting feature names after OneHotEncoder
feature_names = numerical_columns + \
    list(clf.named_steps['preprocessor'].named_transformers_['cat'].
    ↪named_steps['onehot'].get_feature_names_out(categorical_columns))

# Getting feature importances from the model
importances = best_model.named_steps['model'].feature_importances_

# Mapping these importances to the respective feature names
feature_importances = dict(zip(feature_names, importances))

# Sorting the features by importance
```

```

sorted_feature_importances = sorted(feature_importances.items(), key=lambda x:
    ↪x[1], reverse=True)

# Displaying the sorted feature importances
for feature, importance in sorted_feature_importances[:20]:
    print(f"{feature}: {importance}")
import matplotlib.pyplot as plt
import seaborn as sns

# Assuming importances and feature_names are already defined as in the previous
    ↪explanation

# Convert the feature importances to a pandas DataFrame for easier plotting
importance_df = pd.DataFrame({'Feature': feature_names, 'Importance':
    ↪importances})

# Sort the DataFrame by importance
importance_df = importance_df.sort_values(by='Importance', ascending=False)

# Plotting
plt.figure(figsize=(12, 10))
sns.barplot(x='Importance', y='Feature', data=importance_df.head(20))
plt.title('Top 20 Feature Importances for XGB')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.show()

```

```

Census_IsVirtualDevice_0.0: 0.014620620757341385
AVProductStatesIdentifier_53447.0: 0.010690007358789444
AVProductStatesIdentifier_63682.0: 0.00941397063434124
Platform_windows7: 0.009104500524699688
ProductName_win8defender: 0.008944679982960224
AVProductStatesIdentifier_43856.0: 0.005159500986337662
AVProductStatesIdentifier_43927.0: 0.0049601453356444836
OsVer_10.0.0.0: 0.0044375997968018055
Census_FirmwareManufacturerIdentifier_152.0: 0.004132535308599472
AVProductStatesIdentifier_46413.0: 0.0038269313517957926
AVProductStatesIdentifier_44141.0: 0.003764175111427903
AVProductStatesIdentifier_43807.0: 0.003493108320981264
IeVerIdentifier_81.0: 0.0034797738771885633
Census_OEMNameIdentifier_666.0: 0.00345293409191072
EngineVersion_1.1.14901.4: 0.0033428927417844534
Census_OSSkuName_CORE_SINGLELANGUAGE: 0.0033259293995797634
Census_OSInstallLanguageIdentifier_33.0: 0.0032077180221676826
GeoNameIdentifier_192.0: 0.002913816599175334
AppVersion_4.14.17639.18041: 0.0028843588661402464
AVProductStatesIdentifier_51954.0: 0.002883317181840539

```


